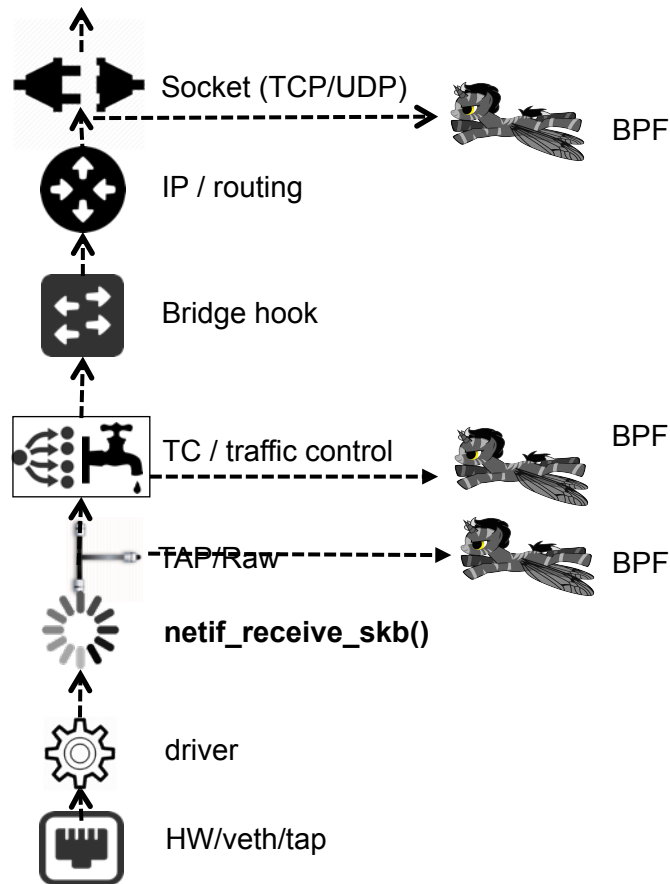


BPF networking examples


Hooking into the Linux networking stack (RX)

- BPF programs can attach to sockets or the traffic control (TC) subsystem
- sockets: TCP, UDP, netlink, af_unix, raw
- This allows to hook at different levels of the Linux networking stack, providing the ability to act on traffic that has or hasn't been processed already by other pieces of the stack




Hooking into the Linux networking stack (TX)

- Opens up the possibility to implement network functions at different layers of the stack




Use BPF to build distributed virtual topology out of existing linux bridges and routers




— = BPF

physical to virtual gateway in BPF

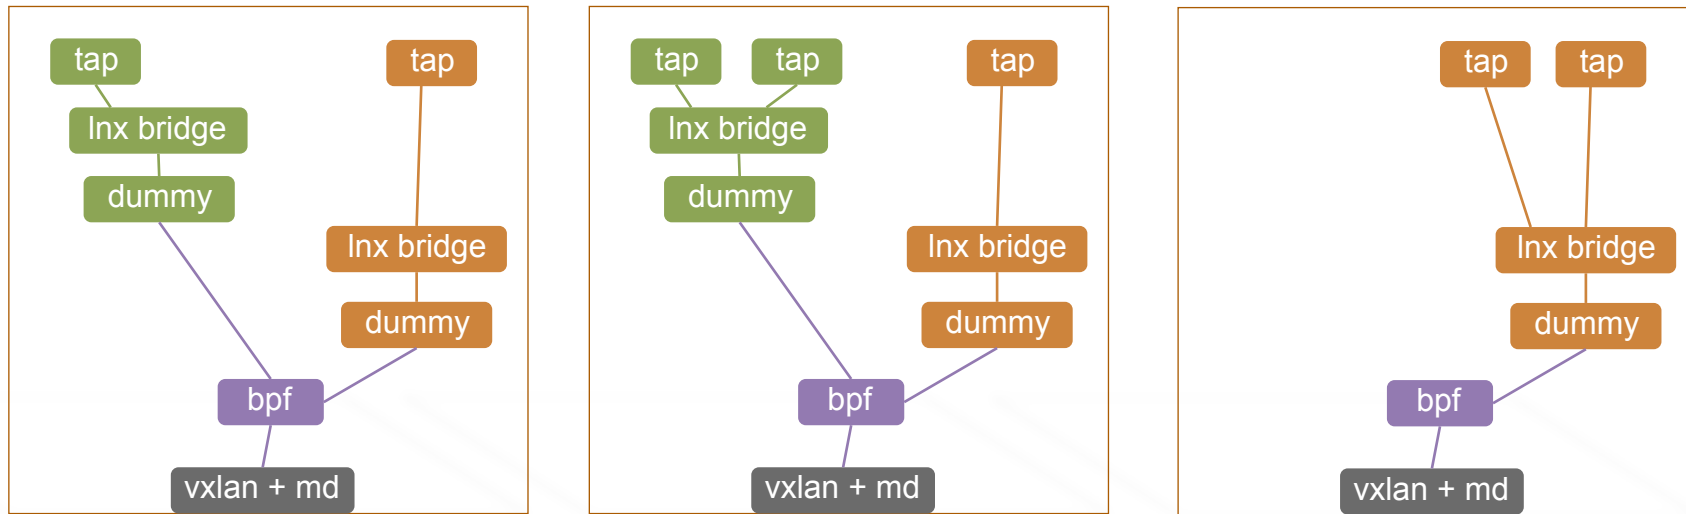


phys2virt: src mac redirect + vlan learning + vlan_pop
virt2phys: ifindex based redirect + vlan_push

Distributed bridge with BPF






Distributed bridge with BPF and multicast



BPF program does

- ingress: redirect {vni} -> {ifindex}, learn {src_mac, ifindex} -> {remote_ip, vni}
- egress: add metadata {dst_mac, ifindex} -> {remote_ip, vni}, redirect to tunnel
use multicast for unknown dst_mac

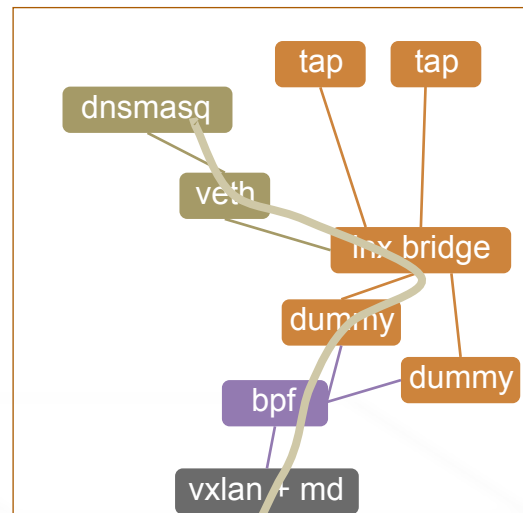


Distributed bridge with BPF (no multicast)



BPF program does

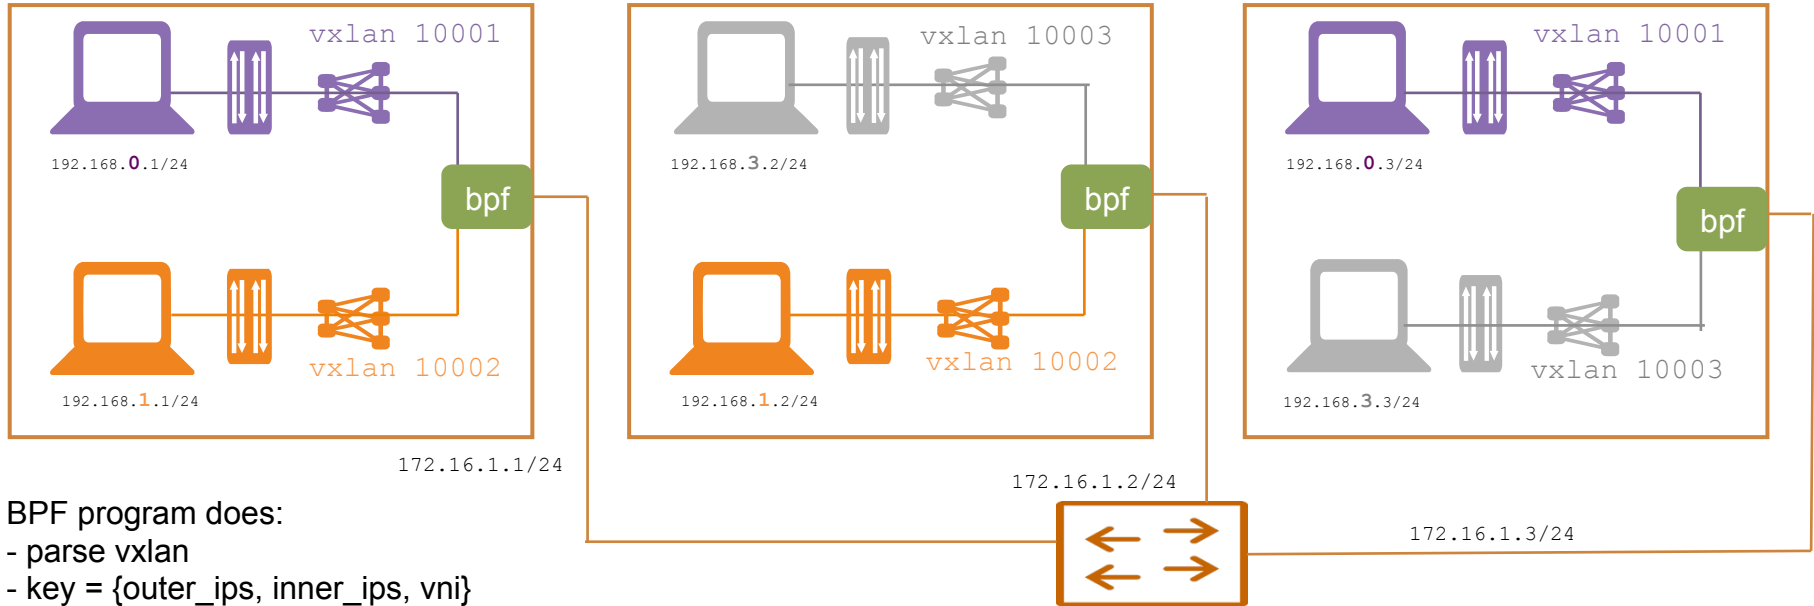
- ingress: redirect {vni, remote_ip} -> {ifindex}
- egress: add metadata {ifindex} -> {remote_ip, vni}, redirect to tunnel

Distributed bridge with BPF (no multicast) and DHCP



packet flow of DHCP request/reply

tunnel analytics with BPF



BPF program does:

- parse vxlan
- key = {outer_ips, inner_ips, vni}
- value = map[key]
- value.rx_pkts ++
- value.rx_bytes += skb->len
- value.tx_pkts ++
- value.tx_bytes += skb->len

tunnel analytics with BPF

- BPF program collects TX/RX stats, stores them in a map

- user space python does:

```
def stats2json(k, v):  
    return {  
        "vni": int(k.vni),  
        "outer_sip": str(IPAddress(k.outer_sip)),  
        "outer_dip": str(IPAddress(k.outer_dip)),  
        "inner_sip": str(IPAddress(k.inner_sip)),  
        "inner_dip": str(IPAddress(k.inner_dip)),  
        "tx_pkts": v.tx_pkts, "tx_bytes": v.tx_bytes,  
        "rx_pkts": v.rx_pkts, "rx_bytes": v.rx_bytes,  
    }  
  
while True:  
    result_total = []  
    for k, v in stats.items():  
        result_total.append(stats2json(k, v))  
    with open("./chord-transitions/data/tunnel.json.new", "w") as f:  
        json.dump(result_total, f)
```

- chord-transitions converts json to chord diagram

