

Abstracting switch architectures - the P4 approach -

Sept 10, 2015

Mihai Budiu

barefootnetworks.com

The P4 tension



Universal

Customizable

P4 v1: Fixed Abstract Forwarding Model

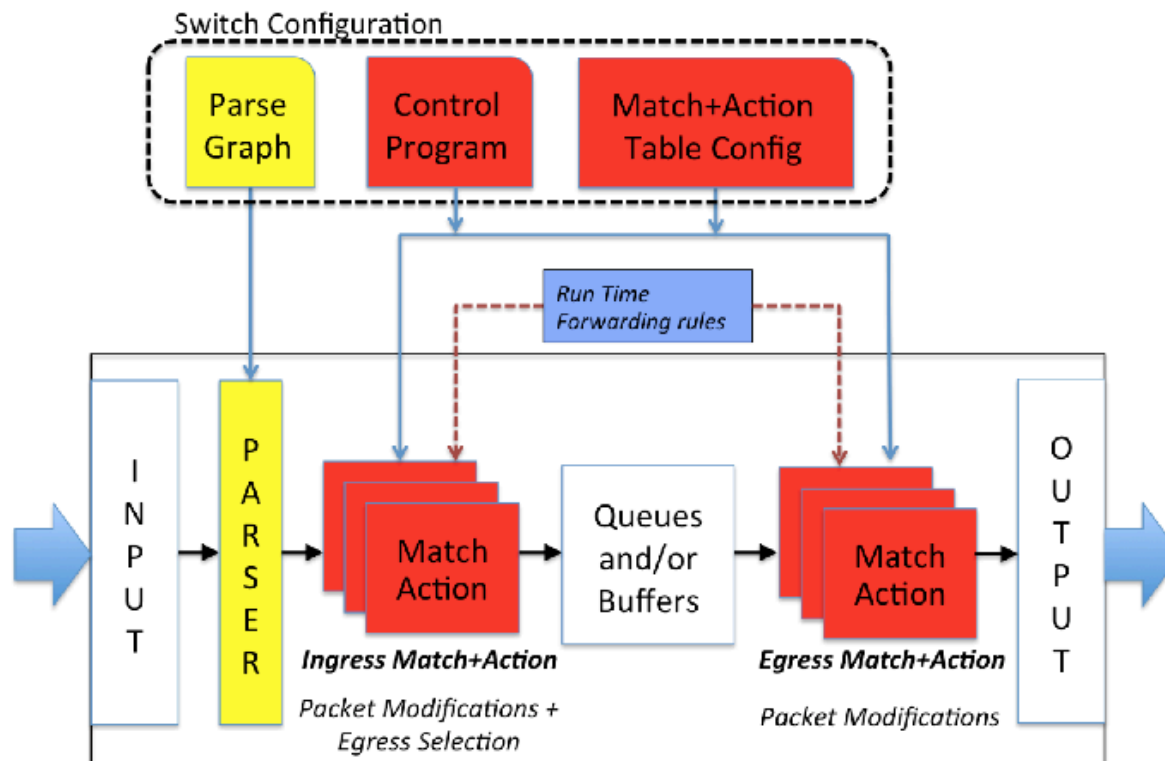




Figure 1: Abstract Forwarding Model

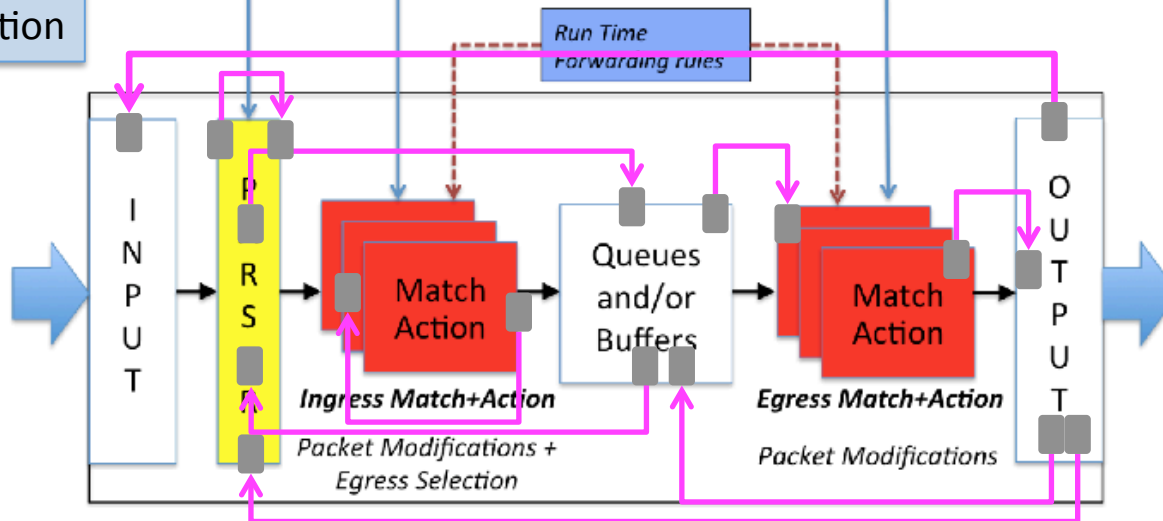

P4 v1: Details

Identifier	Exception Event
p4_pe_index_out_of_bounds	A header stack array in- clared bound.
p4_pe_out_of_packet	There were not plete an
p4_pe_header_too_long	A c st , the declared
p4_pe_header_too_short	was less than the min- xed length portion of the
p4_pe_unh	ment had no default specified but the on value was not in the case list.
	checksum error was detected.
	This is not an exception itself, but allows the programmer to define a handler to specify the default behavior if no handler for the condition exists.

Table 2: Standard Parser Exceptions




Switch Configuration



Field	Notes
ingress_port	The port on which the packet arrived. Set prior to parsing. Always defined. Read only.
packet_length	The number of bytes in the packet. Does not include the CRC. Set by the parser. Cannot be used for matching or the switch is in 'out through'.
egress_spec	Specification of an egress match-action during the 'intended' egress physical port(s) (user physical port, a logical AG, a route, or group).
egress_port	The physical port(s) committed by the egress head.
egress_instance	An instance of a physical port, this is the buffering mechanism match-action stages. See Section 13 below.
instance_type	of instance of the packet: one of the flags or type value; for example, could be 'recirculated'. Do we need a counter tool of the parser. 0 means no error. Otherwise, value indicates what error occurred during parsing. Specific representation is TBD.
parser	A parser error occurred, this is an indication of the location in the parser program where the error occurred. Specific representation is TBD.
parser	

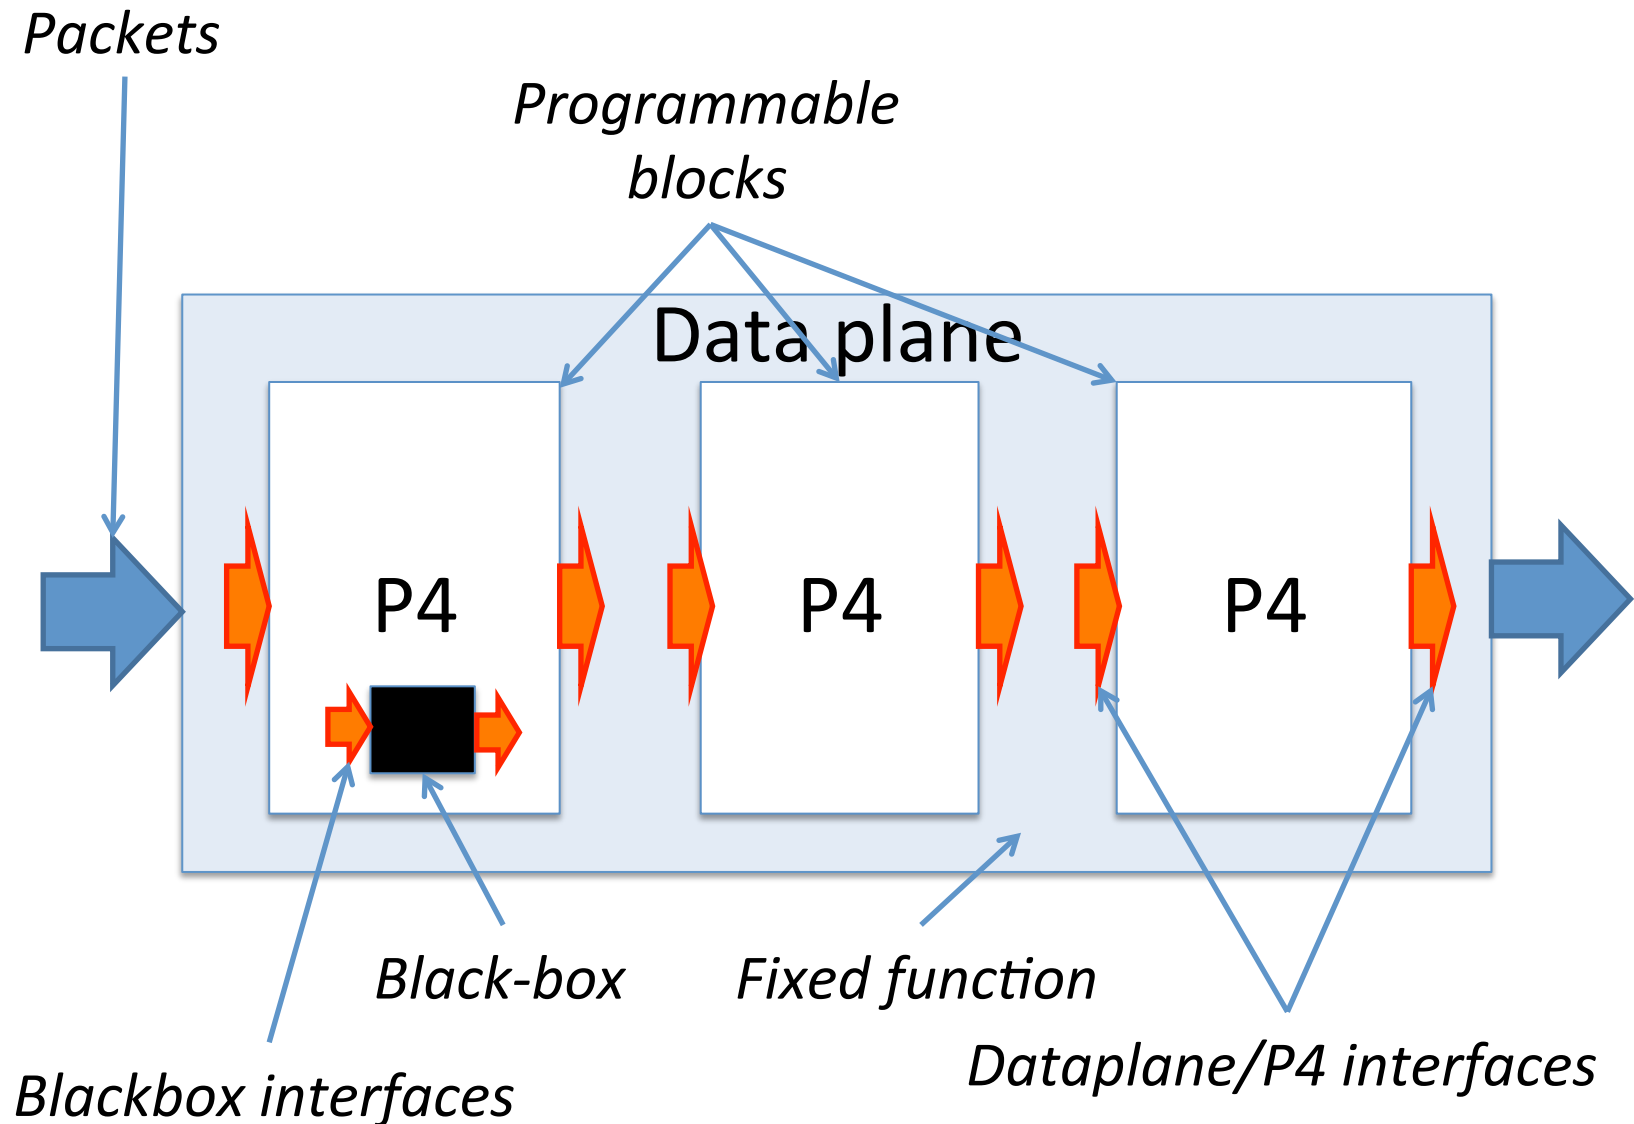
Table 3: Standard Intrinsic Metadata Fields

Divide and conquer

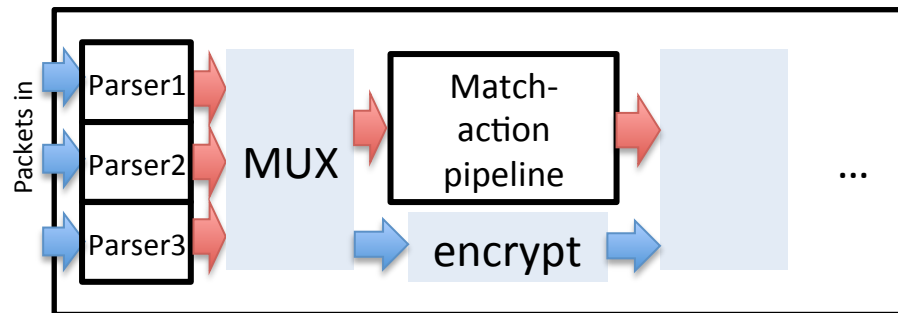

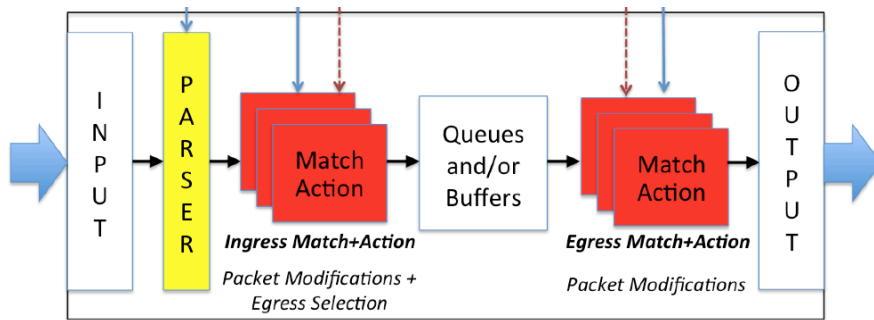


- Separate language definition from architecture definition
- Evolve them independently

Generic Programmable Dataplane Model



P4 Support for multiple architectures



Switch architecture in C++

```
// switch.hpp: written by manufacturer
struct MetaIn { int inputPort; }
struct MetaOut {
    int outputPort;
    bool drop;
}

template<class T> class switch
{
    virtual void parser(const packet &p, T& headers)=0;
    virtual void pipe(T& headers,
        const MetaIn &in, MetaOut& out)=0;
}

target-defined metadata
user-defined metadata
abstract methods = implemented by user
```

The diagram illustrates the switch architecture in C++ with annotations. Red arrows point from the text annotations to specific parts of the code. The annotation 'target-defined metadata' points to the 'MetaOut' struct. The annotation 'user-defined metadata' points to the 'T' template parameter in the 'switch' class. The annotation 'abstract methods = implemented by user' points to the virtual methods 'parser' and 'pipe' in the 'switch' class.

Architecture-specific black-boxes

// target.hpp: written by manufacturer

class Cchecksum

{

Checksum();

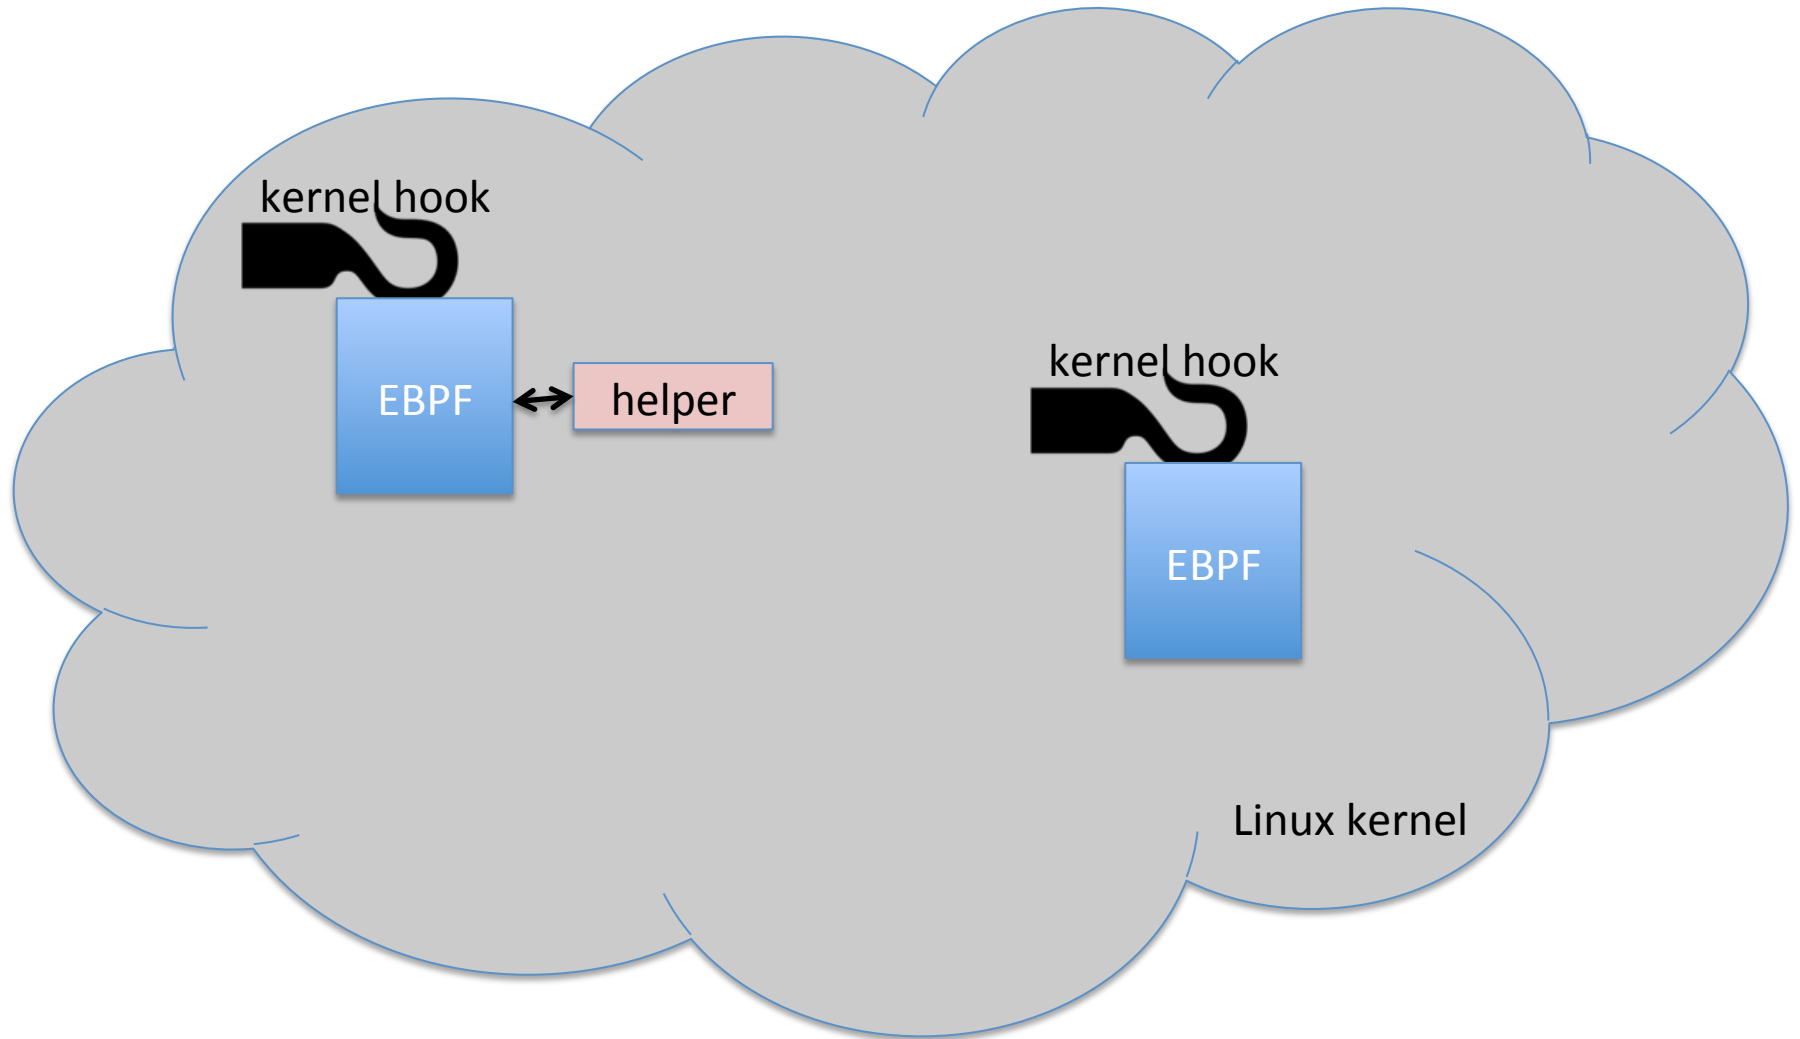
void clear();

template<class T> void append(const T& data);

u16 getChecksum();

}

EBPF's world



Parallels

- P4 switch architecture = EBPF set of hooks
- P4 black boxes = EBPF helper functions
- P4 intrinsic metadata = EBPF hook “signatures”

A suggestion

- Structured interfaces in IOVisor:
 - Header file:
 - List of predefined data types (e.g. skb, TC_*)
 - Prototypes for all EBPF hooks
 - A way to name hooks in a uniform way (TC, socket, network interface, etc.)
 - An abstracted way to attach EBPF code to a hook
 - A list of helper declarations
 - Text file: model description explaining how data flows between hooks